

UNIVERSITY OF WATERLOO
Cheriton School of Computer Science

CS 458/658

Computer Security and Privacy

**Winter 2017
Florian Kerschbaum**

ASSIGNMENT 1

Blog Task signup due date: **Sunday, January 15, 2017 3:00pm (no extension)**

Milestone due date: **Wednesday, January 25, 2017 3:00pm**

Assignment due date: **Friday, February 3, 2017 3:00 pm**

Total Marks: 60 (+5 bonus)

Blog Task TA: Peiyuan Liu

Office hours: Wednesday 12:00–1:00pm, DC 3332

Written Response TA: Mohamed Sabri

Office hours: Wednesday 9:30–10:30am, DC 3314

Programming TA: Yossef Musleh

Office hours: Tuesdays 2:30–3:30pm, DC 3332

Please use Piazza for all communication. Ask a private question if necessary. The TAs' office hours are also posted to Piazza for reference.

Blog Task

0. [0 marks, but -2 if you do not sign up by the due date] Sign up for a blog task timeslot by the due date above. The 48 hour late policy, as described in the course syllabus, does not apply to this signup due date. Look at the blog task in the Course Materials, Content section of the course website to learn how to sign up. Please visit <https://crysp.uwaterloo.ca/courses/cs458/blogtask-scheduling/> for the sign-up page.

Written Response Questions [30 marks]

Note: Please ensure that written questions are answered using complete and grammatically correct sentences. You will be marked on the presentation and clarity of your answers as well as the content.

1. (6 marks total) For each of the following, please:
 - identify the scenario as a compromise of Confidentiality, Integrity, Availability, and/or Privacy,
 - and, briefly explain your choice of compromise.

(Hint: You are only required to give one compromise and corresponding explanation per scenario.)

- (a) (2 marks) An attacker installs malware on your home wireless router and uses it in a distributed denial-of-service-attack.
- (b) (2 marks) You have a regular medical exam. The medical lab transfers your test results to a pharma company to improve a recently developed drug.
- (c) (2 marks) An attacker brute-force guesses your online banking passwords, logs in and views your records.

Most of these are ambiguous. The important thing is that the explanation demonstrates understanding or critical thinking. [1 mark] is given for identifying the compromise. [1 mark] is given for an explanation that matches and supports the chosen compromise.

- (a) (2 marks) An attacker installs malware on your home wireless router and uses it in a distributed denial-of-service-attack.
 - Integrity, since the malware modifies your system to an unintended state.
 - Availability, since the DDOS attacks prevents access to the network.
- (b) (2 marks) You have a regular medical exam. The medical lab transfers your test results to a pharma company to improve a recently developed drug.
 - Privacy, since it is revealing data to someone they didn't intend to give it to.
 - Confidentiality, since medical data is sensitive.
- (c) (2 marks) An attacker brute-force guesses your online banking passwords, logs in and views your accounts.
 - Confidentiality, since the password and banking information is secret information.
 - Privacy, since personal banking data was revealed.
 - Integrity, since the authentication information was forged.

2. (8 marks) Consider an academic conference submission system (such as, for example, HotCRP). Authors can submit their confidential papers, chairs can assign them to blind reviewers who can then read and rate them. Identify four threats that apply to the management of a secure reviewing process, and categorize them as one of the major categories: *interception*, *interruption*, *modification*, and *fabrication*. You must include at least one threat from each category. Be creative.

Furthermore, suggest a potential defence for each of the four attacks, and classify the defence as prevention, deterrence, deflection, detection, or recovery. You may have duplicate types for this part.

Solutions will vary greatly, but will probably be similar to the following. **Marking:** One mark for a correctly classified attack, one for a correctly classified defense.

- **Interception:** A reviewer reading a paper not assigned to them. An author can read papers from other authors, ...
 - **Defences:** Authenticate and control access to papers (prevention)
- **Interruption:** Attacks on the submission system, crashes, ...
 - **Defences:** Harden the operating system (deterrence), keep a backup (recovery),
- **Modification:** Changing a paper (not by the author) or review (not by the reviewer)
 - **Defences:** Append-only, do not overwrite, keep log files (detection), ...
- **Fabrication:** An author performing his own review, ...
 - **Defences:** Chair email notification (detection), use lowest rating (deflection), ...

3. (16 marks total) Cybercrime is still on the rise. Some trends have been identified in a recent article <http://www.techrepublic.com/article/2017-cybercrime-trends-expect-a-fresh-wave-of-ransomware-and-iot-hacks/>.

While the article mentions several important aspects, it also imprecise. The purpose of this question is to explore and critique some of the issues discussed in the article. You will be required to give written answers, using complete sentences formed into well-structured and coherent arguments. Each part should require no more than a couple of paragraphs to answer.

(Hint: You may need other references other than the one listed here. Include your references in your answer)

- (a) (4 marks) The article mentions the number and costs of attacks. How realistic are these numbers?
 - (b) (4 marks) The expert mentions several best practices at the end of the article. Identify at least one important best practice he misses.
 - (c) (4 marks) The article mentions IoT as a threat to privacy. Discuss how that could manifest.
 - (d) (4 marks) Several given best practices are similar – differing only in small aspects. Identify and discuss the similarity of two best practices.
- (a) (4 marks) The number of 90 million attacks is an underestimation, given the fact that the number of connected devices in 2016 was 6.4 billion [1]. And, Symantec reported that the daily number of attacks against individuals alone was more than 1 million [2]. The article's reported cost of \$15 million per attack is also not realistic, it is more likely that this is the average annual cost per company, and not per attack [3].

References:

- [1] <http://www.gartner.com/newsroom/id/3165317>
- [2] <https://www.symantec.com/security-center/threat-report>
- [3] <http://www.ponemon.org/blog/2015-cost-of-cyber-crime-united-states>

- (b) (4 marks) The article misses the following best practices: Backup, encryption, data breach plan, user activity monitoring and logging.
- (c) (4 marks) The advent of IoT presents a great security challenge, as it allows an increasing number of devices to be connected to the Internet. These newly connected devices represent new points of vulnerability; it is much easier for the attacker to breach into a more secure system through a connected device that is less secure. IoT will also allow attacker access to devices and data that were not previously available to them, such as smart door locks and smart in-home appliances. Also, IoT poses concerns for data collection and privacy, since the manufacturing companies can collect data about users and usage patterns.
- (d) (4 marks) The following pairs of suggested best practices are very similar:
- "Restricting user permissions" and "applying the principle of least privilege".
 - "Conducting frequent vulnerability scanning" and "conducting penetration testing".
 - "Patching and updating everything" and "maintaining up-to-date antivirus software".

Programming Question [30 marks + 5 bonus]

Background

You are tasked with testing the security of a custom-developed *file submission application* for your organization. It is known that the application was *very poorly written*, and that in the past, this application had been exploited by some users with the malicious intent of *gaining root privileges*. There is some talk of the application having *three or more vulnerabilities*! As you are the only person in your organization to have a background in computer security, only you can *demonstrate how these vulnerabilities can be exploited* and *document/describe your exploits* so a fix can be made in the future.

Application Description

The application is a very simple program to submit files. It is invoked in the following way:

- `submit <path to file> [message]` : this will copy the file from the current working directory into the submission directory, and append the string “message” to a file called `submit.log` in the user’s home directory.

There may be other ways to invoke the program that you are unaware of. Luckily, you have been provided with the source code of the application, `submit.c`, for further analysis.

The executable `submit` is *setuid root*, meaning that whenever `submit` is executed (even by a normal user), it will have the full privileges of *root* instead of the privileges of the normal user. You can check which user you are running as with the command `whoami`.

Testing Environment

To help with your testing, you have been provided with a virtual *user-mode linux* (uml) environment where you can log in and test your exploits. These are located on one of the *ugster* machines. To obtain your *ugster* account information, go through the quest authentication and visit: <https://crysp.uwaterloo.ca/courses/cs458/blogtask-scheduling/>

Once you have logged into your *ugster* account with SSH, you can use the `uml-a1` command to start your virtual Linux environment. The following logins can be used:

- `user` (no password): main login for virtual environment
- `halt` (no password): halts the virtual environment, and returns you to the *ugster* prompt

The executable `submit` application has been installed to `/usr/local/bin` in the virtual environment, while `/usr/local/src` in the same environment contains `submit.c`. Conveniently, someone seems to have left some shellcode in `shellcode.h` in the same directory.

It is important to note all changes made to the virtual environment will be lost when you halt it. Thus it is important to remember to keep your working files in `/share` on the virtual environment, which maps to `~/uml/share` on the ugster environment.

Rules for exploit execution

- You have to submit three exploit programs to be considered for full credit. One of your exploit programs **MUST** target either a buffer overflow vulnerability or a format string vulnerability. The other two may target other vulnerabilities.
- Each vulnerability can be exploited only in a single exploit program. A single exploit program can exploit more than one vulnerability. If unsure whether two vulnerabilities are different, please ask a private question on Piazza.
- There is a specific execution procedure for your exploit programs (“*sploits*”) when they are tested (i.e. graded) in the virtual environment:
 - Sploits will be run in a **pristine** virtual environment, i.e. you should not expect the presence of any additional files that are not already available
 - Execution will be from a clean `/share` directory on the virtual environment as follows: `./sploitX` (where `X=1..3`)
 - Sploits must not require any command line parameters
 - Sploits must not expect any user input
 - If your sploit requires additional files, it has to create them itself
- For marking, we will compile your exploit programs in the `/share` directory in a virtual machine in the following way: `gcc -Wall -ggdb sploitX.c -o sploitX`. You can assume that `shellcode.h` is available in the `/share` directory.
- Be polite. After ending up in a root shell, the user invoking your exploit program must still be able to exit the shell, log out, and terminate the virtual machine by logging in as user `halt`. None of the exploits should take more than about a minute to finish.
- Give feedback. In case your exploit program might not succeed instantly, keep the user informed of what is going on.

Deliverables

Each sploit is worth 10 marks, divided up as follows:

- 6 marks for a successfully running exploit that gains root
- 4 marks for a description of the vulnerability used, an explanation of how your sploit program exploits the vulnerability, and a description of how the vulnerability could be fixed

A total of three exploits must be submitted to be considered for full credit, with at least one being a *buffer overflow* or *format string* exploit. Marks may be docked if you do not submit a buffer overflow or format string exploit.

Bonus exploits

You may submit *at most one* extra vulnerability for bonus points. It must exploit a distinct vulnerability from your first three exploits. Bonus marks will be given as follows:

- 3 marks for a successfully running exploit that gains root
- 2 marks for a description, as above. Note that bonus points *will not be awarded* for descriptions without accompanying working exploit code.

There are at least five vulnerabilities in the program:

1. Buffer overflow in the print usage function. It copies up to 640 characters into a 281 character array. This can be exploited by changing the first argument in the `execv` call to a NOP sled + shellcode + ret addr.
2. There is another buffer overflow in the copy file function. This takes a little bit more work than a simple NOP sled + shellcode + retaddr in the submitted file. The virus checker will catch any references to `/bin/sh`. Instead, make another program (called `/tmp/fo` or something else) and have it call `/bin/sh` as root.
3. There is a format string vulnerability in the print version function. This can be exploited by a formatting string overflowing the buffer by adding a print out.
4. There is a `tocttou` vulnerability in the amount of time it takes to do the virus check/verifying the log file path is safe.
5. There is a simple vulnerability where you can change the environment path, and have it point to a new 'mkdir' command that spawns a shell.

What to hand in

All assignment submission takes place on the `student.cs` machines (not `ugster` or the virtual environments), using the `submit` utility. In particular, log in to the Linux student environment (`linux.student.cs.uwaterloo.ca`), go to the directory that contains your solution, and submit using the following command: `submit cs458 1 .` (dot included). CS 658 students should also use this command and ignore the warning message.

By the **milestone due date**, you are required to hand in:

spl0it1.c One completed exploit programs for the programming question. Note that we will build your spl0it programs **on the uml virtual machine**

a1-milestone.pdf: A PDF file containing your answers for the written-response questions, and the exploit descriptions for spl0it1.

Note: You will not be able to submit spl0it1.c, or a1-milestone.pdf after the milestone due date (plus 48 hours).

By the **assignment due date**, you are required to hand in:

spl0it2.c, spl0it3.c: The two remaining exploit programs for the programming question.

(optional) spl0it4.c: Exploit program for the programming bonus question.

a1.pdf: A PDF file containing your exploit descriptions for spl0it2, spl0it3 and optional spl0it4. Do not put written answers pertaining to spl0it1 and written-response questions into this file; they will be ignored.

The 48 hour late policy, as described in the course syllabus, applies to the milestone due date and the assignment due date. It does not apply to the blog task signup due date.

Useful Information For Programming Sploits

Most of the exploit programs do not require much code to be written. Nonetheless, we advise you to start early since you will likely have to read additional information to acquire the necessary knowledge for finding and exploiting a vulnerability. Namely, we suggest that you take a closer look at the following items:

- Module 2
- Smashing the Stack for Fun and Profit (<http://insecure.org/stf/smashstack.html>)
- Exploiting Format String Vulnerabilities (v1.2) (<http://julianor.tripod.com/bc/formatstring-1.2.pdf>) (Sections 1–3 only)
- The manpages for `execve` (`man execve`), `pipe` (`man pipe`), `popen` (`man popen`), `getenv` (`man getenv`), `setenv` (`man setenv`), `passwd` (`man 5 passwd`), `shadow` (`man 5 shadow`), `symlink` (`man symlink`), `expect` (`man expect`).

GDB

The `gdb` debugger will be useful for writing some of the exploit programs. It is available in the virtual machine. In case you have never used `gdb`, you are encouraged to look at a tutorial (e.g., <http://www.unknownroad.com/rtfm/gdbtut/>).

Assuming your exploit program invokes the `submit` application using the `execve()` (or a similar) function, the following statements will allow you to debug the `submit` application:

1. `gdb sploitX` (`X=1..3`)
2. `catch exec` (This will make the debugger stop as soon as the `execve()` function is reached)
3. `run` (Run the exploit program, which will stop when the exec of `submit` happens)
4. `symbol-file /usr/local/bin/submit` (We are now in the `submit` application, so we need to load its symbol table)
5. `break main` (Set a breakpoint in the `submit` application)
6. `cont` (Run to breakpoint)

You can store commands 2–6 in a file and use the “`source`” command to execute them. Some other useful `gdb` commands are:

- “`info frame`” displays information about the current stackframe. Namely, “`saved eip`” gives you the current return address, as stored on the stack. Under saved registers, `eip` tells you where on the stack the return address is stored.
- “`info reg esp`” gives you the current value of the stack pointer.

- “x <address>” can be used to examine a memory location.
- “print <variable>” and “print &<variable>” will give you the value and address of a variable, respectively.
- See one of the various gdb cheat sheets (e.g., <http://darkdust.net/files/GDB%20Cheat%20Sheet.pdf>) for the various formatting options for the print and x command and for other commands.

Note that `submit` will not run with root privileges while you are debugging it with gdb. (Think about why this limitation exists.)

The Ugster Course Computing Environment

In order to responsibly let students learn about security flaws that can be exploited in order to become “root”, we have set up a virtual “user-mode linux” (uml) environment where you can log in and mount your attacks. The gcc version for this environment is the same as described in the article “Smashing the Stack for Fun and Profit”; we have also disabled the stack randomization feature of the 2.6 Linux kernel so as to make your life easier. (But if you’d like an extra challenge, ask us how to turn it back on!)

To access this system, you will need to use ssh to log into your account on one of the ugster machines: `ugsterXX.student.cs.uwaterloo.ca`. There are a number of ugster machines, and each student will have an account for one of these machines. To obtain your ugster account information, visit: <https://crysp.uwaterloo.ca/courses/cs458/blogtask-scheduling/> under *View Grades*. Authentication is done via UW CAS and uses your quest student information.

The ugster machines are located behind the university’s firewall. While on campus you should be able to ssh directly to your ugster machine. When off campus, you have the option of using the university’s VPN (see these instructions), or you can first ssh into `linux.student.cs.uwaterloo.ca` and then ssh into your ugster machine from there.

When logged into your ugster account, you can run “`uml-a1`” to start the user-mode linux to boot up a virtual machine.

The gcc compiler installed in the uml environment may be very old and may not fully implement the ANSI C99 standard. You might need to declare variables at the beginning of a function, before any other code. You may also be unable to use single-line comments (“//”). If you encounter compile errors, check for these cases before asking on Piazza.

Any changes that you make in the uml environment are lost when you exit (or upon a crash of user-mode linux). **Lost Forever.** Anything you want to keep must be put in `/share` in the virtual

machine. This directory maps to `~/uml/share` on the ugster machines, which is how you can copy files in and out of the virtual machine. It can be helpful to ssh twice into the ugster. In one shell, start user-mode linux, and compile and execute your exploits. In the other shell, edit your files directly in `~/uml/share/`, to ensure you do not lose any work. The ugster machines are not backed up. You should copy all your work over to your student.cs account regularly.

When you want to exit the virtual machine, use `exit`. Then at the login prompt, login as user “halt” and no password to halt the machine.

Any questions about the ugster environment should be directed to the Programming Question TA.